

UI Review & Product Logic Response

This report responds to a 7-page UI / product logic review. Each page compares the original red-dot feedback with the updated interface, explains how SpeakSpark now feels more like a private practice coach, and separately calls out newly added Navigator guidance, the Trying Point loop, and the calendar history view.

Report date **2026-06-19** Source review **speakspark-labelled-design-walkthrough.pdf**

Positioning **Private practice coach, not a scoring tool**

● Shipped

● Partial / Todo


● Backend-dependent (restored; AI working)

PAGE 2


Start Screen — Motivation Without Scoring

The first screen is friendly, but scores and streaks can create practice pressure.

VERSION 1 REVIEWER ORIGINAL



VERSION 2 CURRENT VERSION



Home V2 (trying point + modes + daily starter
+ no streak/score) · #/

OUR RESPONSE (Numbers match the red dots above)

1 Streak creates pressure ● Shipped

Original suggestion: use "5 practices this week" instead of a streak.

✓ Changed to "5 practices this week", removing streak pressure.

87a6cd6

2 "92 pts" feels like a score ● Shipped

Original suggestion: replace with "2 phrases saved / 1 coach note".

✓ Removed the raw score and changed it to "Phrase saved · ... win" as evidence.

87a6cd6

3 CTA leans toward competition ● Shipped

Original suggestion: test "Start Practice".

✓ Changed to "Start Practice" (Home + Practice).

87a6cd6

+ Added: Home Navigator guidance ● Shipped


New users see a 3-step guide (mode / topic / Start), shown only once.

63b78d4


Practice Screen — Reduce Judging Signals

While the student is speaking, the interface should feel like a calm coach, not a live scorecard.

VERSION 1 REVIEWER ORIGINAL



VERSION 2 CURRENT VERSION



Practice V2 (muted unscored KTV metrics + trust label) · [#/practice](#)

OUR RESPONSE (Numbers match the red dots above)

1 "AI moments" feels vague and score-like

● Shipped

Original suggestion: change to "coach notes" or remove.

✓ Removed the vague count and replaced it with the top 🎯 topic / mode.

f6392c9

2 Metrics feel like live judgment

● Shipped

Original suggestion: hide metrics during speaking, then use them as private evidence after the run.

✓ Reframed them to serve as an encouragement frame: unscored metrics are muted, scored ones gain color, and fresh gains highlight so attention follows progress.

87a6cd6

3 "AI is listening" feels like surveillance

● Shipped

Original suggestion: change to "Coach is following your idea".

✓ Changed to "Coach is following your story" + the cat state.

fa5b2c2

35472de

4 Loading blocks practice

● Shipped


Original suggestion: allow recording first, generate later.

✓ Added pre-generated buffers and layered fallbacks; pause cards prefer real AI (instant when buffered; otherwise the cat thinks; local fallback only when offline). Backend is restored and real AI is working.

1936038

Takeaway Top — Put The Next Move Before The Trophy

The first screen after a run should be a coach plan, not a score sheet.



OUR RESPONSE (Numbers match the red dots above)

1 Trophy implies winning or ranking

● Shipped

Original suggestion: notebook/check + "Practice complete".

✓ Changed to "Practice complete" + 🎉, with no trophy or ranking.

838a68b

2 Stats feel like a report card

● Shipped

Original suggestion: start with a human sentence before data.

✓ Start with a human headline ("You kept your idea going...").

838a68b

3 Next Run Plan should be the hero

● Shipped

Original suggestion: promote it into the top hero card.

✓ Next Run Plan now appears in the first-screen Level up, so it no longer feels like an afterthought.

fd71e59

4 Red AI failure looks like an error

● Shipped

Original suggestion: use neutral copy.

✓ Changed to neutral, honest amber notice (this is what helped locate the backend failure).

1936038

ca2982f



Added: Takeaway V2 information architecture

● Shipped

The end page is reordered as ① concise encouragement → ② Level up → ③ details → ④ Coach. Level up combines Amplify + Change/Next time and directly tells the student what to change, how to say it, which words to reuse, and what the next trying point is.

fd71e59

6e889dd



Added: concrete feedback from real AI

● Shipped

Takeaway is generated from the actual transcript. It no longer talks about the lowest score; it outputs specific sentence frames, vocabulary upgrades, and reusable expressions.

1936038


ef98cbf

6e889dd


Takeaway Lower — Guard The Coach Chatbox

The chatbox is powerful, so its boundaries need to be visible before launch.

VERSION 1 REVIEWER ORIGINAL



VERSION 2 CURRENT VERSION



Block ④ Coach (**real AI reply**): based on this run's opening line + presets such as "Ask me one question" (to avoid writing answers for the student) + Practice Again.

OUR RESPONSE (Numbers match the red dots above)

1 Prompts can easily generate full answers

● Shipped

Original suggestion: "Go deeper" → "Ask me one question".

✓ Presets now use "Ask me one question" and similar prompts, so the coach does not answer for the student.

fa5b2c2

2 Progress map falls back into scores

● Shipped

Original suggestion: use text trends; make numbers secondary or hidden.

✓ Changed to **text trends** (Strong / Growing / Just starting + growth descriptions).

4bdaf81


3 Practice Again CTA is strong

● Kept

✓ Kept the core loop and added a persistent coach-mode trust label.

Growth Page Top — Keep The Trying Point, Rename The Scores

The growth page structure is right, but "Ability Portrait" still feels evaluative.



OUR RESPONSE (Numbers match the red dots above)

1 Growth map structure is right

● Kept

- ✓ Kept the personal growth language and avoided public ranking.

2 "Trying Point" works well

● Kept

Original suggestion: keep and strengthen this personal-growth language.

- ✓ My keeps Trying Point on the first screen as a visible goal for the next run.

3 "Ability Portrait" feels like a scorecard

● Shipped

Original suggestion: rename it "Practice growth".

- ✓ Renamed to "Practice growth".

87a6cd6


+ Added: cross-page Trying Point loop

● Shipped

Added a shared source: this week's trying point card on Home (tap into Practice) + My list + Takeaway's "Your next trying point", so the next-run goal becomes a loop.

Growth Page Lower — Make Coach Mode A Trust Feature

Settings should make the content-safety boundary clear to parents, teachers, and students.



OUR RESPONSE (Numbers match the red dots above)

- 1 Practice history is useful evidence** ● Kept
- ✓ Kept Practice history so past sessions become the student's own growth evidence, not a public ranking or score wall.

- 2 "User management" is too internal** ● Shipped
- Original suggestion: rename it "Settings".
- ✓ Renamed to "**Settings**".
- 87a6cd6

- 3 Coach mode should be clearer** ● Shipped
- Original suggestion: "Gentle prompts" → "Practice feedback only".
- ✓ Changed to "**Practice feedback only**"; the trust label was also extended to Practice and the My header.

- + Added: Calendar history view** ● Shipped
- The calendar opens from the top right; practice days have blue dots, and selecting a day filters the fixed-height list below to that day's sessions so the page no longer grows endlessly.


87a6cd6

6e889dd

This update



Added: real Settings



Settings is no longer a placeholder: Coach mode is an information row, Saved phrase bank reads real historical phrases, and Coaching prompt can be edited and injected into the AI coach.

Appendix Five Additional Text Feedback Items (#1-#5)

Additional trial feedback outside the review PDF, mostly about Practice and Home, has been implemented alongside the page updates above.

#1 Single-line Karaoke captions + fixed End button

● Shipped

ORIGINAL ISSUE

Multi-line captions overloaded the screen; in noisy environments the End button could be pushed out of reach.

59156f9

9995d42

9e8cb75

WHAT WE CHANGED

Teleprompter-style captions: current sentence locked in the middle, older text fades upward, fixed height; End stays fixed in the top bar.

#2 Removed redundant waveform motion

● Shipped

ORIGINAL ISSUE

Waveform + captions + bubbles + KTV + cards created too many moving elements.

9e8cb75

f6392c9

fc53fd2

WHAT WE CHANGED

Removed the 40-bar waveform animation; the cat is just the cat; repeated listening copy is consolidated into one fixed status slot.

#3 Topic reminder + cold-start keywords

● Shipped

ORIGINAL ISSUE

Students can forget the topic at cold start or lose the thread, especially on unfamiliar science topics.

dee6b0f

f6392c9

WHAT WE CHANGED

Persistent top 🌀 topic / mode; cold-start keyword chips turn green when mentioned.

#4 Top-down Takeaway + transcript quotes

● Shipped (real AI working)

ORIGINAL ISSUE

Advice felt disconnected from what was actually said; users could not see "because I said X, I should try Y".

fa5b2c2

838a68b

4bdaf81

6e889dd

WHAT WE CHANGED

Start with "what you said" before advice; each item includes a "You said: ..." quote. The page was rebuilt into three major long-form blocks; the latest version merges Change + Next time and requires concrete frames, words, and examples without mentioning the lowest score.

#5

Themed speaking activity on Home

● Merged (deduplicated)

ORIGINAL ISSUE

Home felt like a dashboard and lacked guidance on what to practice next.

WHAT WE CHANGED

Use **practice mode selection** (Free Talk/Exam/Story) + **today's topic starter** to answer "what should I practice?" The previously added standalone themed activity card duplicated those two controls, so it was **removed to avoid redundancy**. A separate theme entry can be added later if needed.

9410ca1

87a6cd6

c V2 Feature Set at a Glance

Features shipped and demo-ready in this phase (backend restored; AI outputs real content).

⚠️ Honest experience note: pause cards still have slight latency

Real AI is now available online, but when the buffer misses or the network is slow, the follow-up card after a pause can still feel half a beat late. The current version uses pre-generation plus local fallback to avoid interruption; V3 will continue improving response speed.

- ✓ Practice teleprompter captions: current sentence locked center, older text fades upward, fixed height
- ✓ Cat coach + one fixed status slot (removed waveform / redundant motion)
- ✓ Pause response (pre-generated buffer, real AI first, local fallback; online still has a slight wait when the buffer misses)
- ✓ New Takeaway order: concise encouragement → Level up → Details → Coach
- ✓ Advice includes transcript quotes ("You said: ...") + honest fallback when AI fails
- ✓ Navigator first-use guidance (Home + two Practice tours)
- ✓ Trying Point loop (Home → Takeaway → My)
- ✓ Coach-mode trust label "Practice feedback only" extended across the app
- ✓ Home Last Highlight uses real history + first-time empty state (no hardcoded fake data)
- ✓ Friendly, actionable fallback copy for mic off / unsupported browser / offline states
- ✓ Topic bar + cold-start keywords (turn green when spoken)
- ✓ KTV four metrics: unscored states muted, scored states colored, fresh gains highlighted
- ✓ Captured words fixed to one horizontal scrolling row; new chips enter from the right and older ones slide out
- ✓ Takeaway merges Change + Next time: one improvement point directly connects to an upgraded sentence frame, reusable words, and trying point
- ✓ Home: removes streak / score pressure, adds trying point and daily topic
- ✓ My: renamed Practice growth / Settings, real phrase bank, fixed-height history list + calendar date view
- ✓ Editable Coaching prompt (injected into AI) + history by calendar date
- ✓ Bottom navigation across four pages (including the missing Home bottom nav fix)
- ✓ Practice pre-start faint placeholder frames (coach card + Transcript, so the page does not feel empty)
- ✓ AI latency: real-time calls disable thinking for speed (Takeaway keeps thinking for quality; pause cards still need optimization)

Backend and Storage Plan + Decisions

Current state:**No login system**("Alex" is a hardcoded mock); history, phrases, and prompt preferences only live in**local localStorage**(lost on device switch or cache clear, and not tied to a user). To support user binding, cross-device continuity, and audio replay, backend storage is needed.

What to store (by user / by session)

DATA	NOW	BACKEND NEEDED?
User account (who, admin)	mock "Alex"	Requires Auth login - the anchor for everything
Practice session(time / mode / duration / score)	localStorage, single device	Requires sessions table (by user)
Transcript(what the student said)	Stored inside the local session	Store with the session in backend
Audio recording	Not recorded or stored	Requires recording + Supabase Storage (audio bucket), linked to session
Coach history(AI Q&A / takeaway from each run)	Regenerated each time; not saved	Store by session to enable review
Phrase bank / trying point / prompt preferences	Derived / stored locally	Store by user (cross-device)
Frontend UI / interaction	Completed	No backend needed

How user admin and multiple sessions split

Everything is rooted in `user_id`(from login): each user only sees their own data (row-level security / RLS). Each practice run = one `sessions` row (`id · user_id · time · mode · duration · score · transcript · audio_path`), with child `highlights`、`coach_messages`(foreign key: `session_id`). "User admin" means the missing login/account layer.

Decision items (team to decide)

- ☐ Should we implement real login / accounts (auth) now?This is required for user binding, cross-device use, and audio. If not, keep single-device localStorage (enough for demo).

- ☐ **Should we record and store audio?**For 11-15 year-old minors, this involves privacy, consent, and storage cost. Alternatively, store only transcript first and skip replay.
- ☐ **Which backend should we use:**restore the old Supabase project or create a new one? Requires Postgres tables (sessions/highlights/coach_messages) + Storage bucket + RLS.
- ☐ **Data retention / privacy:**Retention period, consent, and deletion mechanism for minors' data, especially audio.
- ☐ **Quota / cost:**Gemini call and storage quota, who pays, and limit strategy.

☒ **Backend status: restored**

Supabase backend is restored, and AI (real-time coach / takeaway summary / Coach Q&A)**is working and demo-ready**. The storage / account / recording items above are decisions about whether to go further into user binding and cross-device use; they are separate from current availability. See the separate **User Guidance**.

E Next Steps / V3 Roadmap

Ongoing iteration, ordered by priority.

- ☐ **Real account system (auth) + database persistence:** Supabase Auth → `user_id` + RLS; sessions/highlights/coach_messages into the database (replacing `localStorage`). This is the prerequisite for cross-device use and user binding.
- ☐ **Recording + minors compliance:** decide whether to record/store audio (Storage), plus consent / retention / deletion before building.
- ☐ **Mobile speech recognition:** browser STT does not support mobile Safari → use server-side STT (repo already has `analyze-voice` as a starting point) or integrate a third party (Deepgram / Whisper).
- ☐ **KTV scoring calibration:** currently local heuristics (IELTS/TOEFL dimensions are directionally right, scores are uncalibrated) → calibrate with real recordings or switch to AI scoring.
- ☐ **Deploy gemini-proxy:** make "disable thinking for real-time calls" effective in production (code is ready).
- ☐ **Pause-card latency optimization:** keep improving pre-generation hit rate, shorten prompts, parallelize requests, or stream responses so follow-up prompts appear faster after a pause.
- ☐ **Analytics / instrumentation:** record behavior events in backend to understand funnel/retention; use real evidence for positive feedback instead of pressure-based streaks.